

STA 221: LECTURE 15

KRISHNA BALASUBRAMANIAN


(UNIVERSITY OF CALIFORNIA, DAVIS)

OUTLINE

- ▷ Introduction to computer architecture
- ▷ Multi-core computing, distributed computing
- ▷ Multi-core computing tools
- ▷ Distributed computing tools

COMPUTER ARCHITECTURE


- ▷ Each computer can have multiple CPUs, each CPU has multiple cores
(e.g., two quad-core CPUs)
- ▷ All the CPUs are connected to memory (e.g., 64G memory)
- ▷ CPU cores can execute in parallel



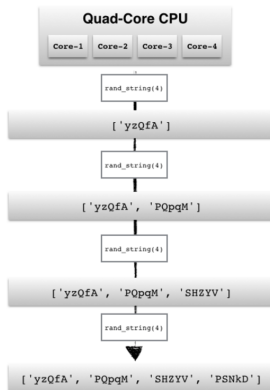
MULTICORE PROGRAMMING

- ▷ Execute tasks simultaneously on many CPU cores

[parallel processing]

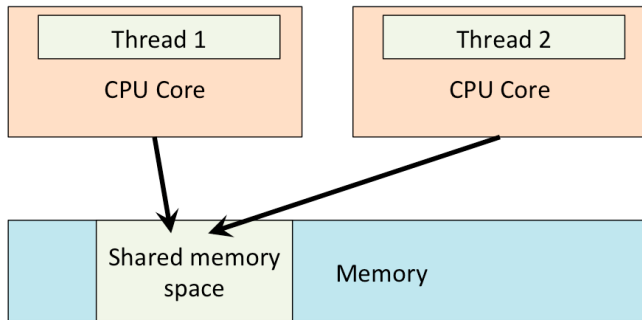


[serial processing]



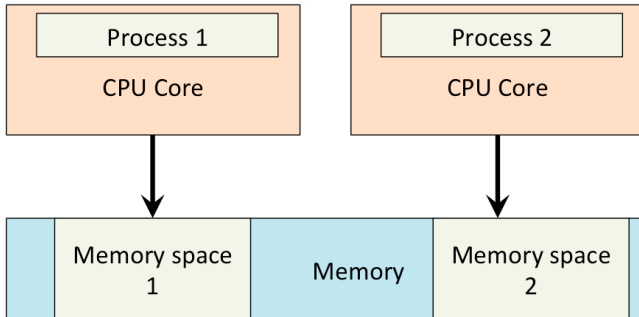
WHAT IS A THREAD?

- ▷ Multiple threads share the memory.
- ▷ Don't need inter-process communication.
- ▷ They are “light-weighted” (not much overhead to fork multiple threads)



WHAT IS A PROCESS?

- ▷ Processes are “share nothing”
(independent executing without sharing memory or state)
- ▷ Easier to turn into a distributed application.



- ▷ Package “threading”
- ▷ Unfortunately, python only allows a single thread to be executing at once
(due to GIL (global interpreter lock))
- ▷ Usually no or little speedup
only useful when you want to interleave I/O and CPU execution

- ▷ Package “multiprocessing”
- ▷ You can create multiple processes
 - ▷ Automatically run on multiple CPU cores
 - ▷ Default no shared memory, each process has its own memory space (larger memory overhead)
 - ▷ Can declare some part of memory to be shared (but often harder to use)
- ▷ You can also check some tutorials:
 - ▷ http://sebastianraschka.com/Articles/2014_multiprocessing.html
 - ▷ <https://pymotw.com/2/multiprocessing/basics.html>

EXAMPLE: HELLOWORLD

```
import multiprocessing as mp

def helloworld(x):
    print 'Hello World %d\n'%x

# Setup a list of processes
plist = []
for x in range(4):
    plist.append(mp.Process(target=helloworld, args=(x,)))

# Run processes
for p in plist:
    p.start()

# Exit the completed processes
for p in plist:
    p.join()
```

EXAMPLE: HELLOWORLD (OUTPUT)

Output of the program:

Hello World 0

Hello World 1

Hello World 2

Hello World 3

BASIC FUNCTIONS

- ▷ (Check <https://docs.python.org/2/library/multiprocessing.html>)
- ▷ “Process(target=helloworld, args=(x,))”:
 - ▷ Specify the target function to run (helloworld)
 - ▷ Specify the input argument of the function (only one argument x)
 - ▷ Create an object belongs to “Process” type
- ▷ The process will run when execute “process.start()”
- ▷ The process will terminate when execute “process.join()”

EXAMPLE: EXCHANGING OBJECTS USING QUEUE

```
import multiprocessing as mp

def f(x,q):
    q.put(x**2)
    return

q = mp.Queue()
processes = []
for x in range(4):
    processes.append(mp.Process(target=f, args=(x,q)))

for p in processes:
    p.start()
for p in processes:
    p.join()

while (q.empty==False):
    print q.get()
```


EXAMPLE: EXCHANGING OBJECTS USING QUEUE (OUTPUT)

Output of the program:

0
1
4
9

USE OF QUEUE

- ▷ `mp.Queue` is a concurrent and “first in first out” data structure
- ▷ Can be used to communicate, or gather the results from the processes
- ▷ `Queue.put()`: insert an object to the end of queue
- ▷ `Queue.get()`: remove the first element in the queue
- ▷ `Queue.empty`: check whether the queue is empty



USING POOL

- ▷ Pool class is another and more convenient approach for parallel processing in python.
- ▷ Use “`mp.Pool(processes=4)`” to create 4 processes
- ▷ Use “`[r1, r2, ..., rk] = pool.map(f, [x1, x2, ..., xk])`” to run multiple processes and get the results
 - ▷ f is the function to run for the processes
 - ▷ $[x_1, \dots, x_k]$ are the input arguments we want to run for the function (this is a size k list)
 - ▷ $[r_1, \dots, r_k]$ are the output arguments we get after running the functions for each input (this is a size k list)
 - ▷ k may be larger than number of processes

EXAMPLE: USING POOL

```
import multiprocessing as mp

def f(x):
    return x**2

pool = mp.Pool(processes=4)
results = pool.map(f, range(4))
print results
```

Output of the program:

```
[0, 1, 4, 9]
```


EXAMPLE: COMPUTE SUM OF SQUARE

```
import multiprocessing as mp

def f(x):
    return x**2

pool = mp.Pool(processes=4)
results = pool.map(f, range(100))
print sum(results)
```

Output of the program:

328350

EXAMPLE: PARALLEL KERNEL DENSITY ESTIMATION

Parzen-window kernel density estimator:

$$p(x) = \frac{1}{n} \sum_{i=1}^n \frac{1}{h^2} \phi\left(\frac{x_i - x}{h}\right)$$


Parallel Implementation in Lecture15.ipynb

Parallel algorithms can be different in the following two cases:

- ▷ Shared Memory Model (Multiple cores)
 - ▷ Independent L1 cache
 - ▷ Shared/independent L2 cache
 - ▷ Shared memory
- ▷ Distributed Memory Model
 - ▷ Multiple processes in single machine
 - ▷ Multiple computers

SHARED MEMORY MODEL (MULTIPLE CORES)

- ▷ Shared memory model: each CPU can access the same memory space
- ▷ Programming tools:
 - ▷ C/C++: openMP, C++ thread, pthread, intel TBB, ...

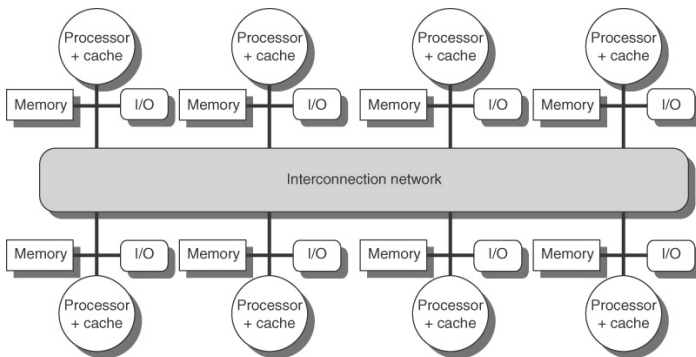


PARALLEL FOR LOOP IN OPENMP

```
#pragma omp parallel for private(i)
for(i=0;i<w_size;i++)
    g[i] = w[i] + g[i];
```

DISTRIBUTED MEMORY MODEL

- ▷ Programming tools: MPI, Hadoop, Spark, ...




© 2007 Elsevier, Inc. All rights reserved.

(Figure from <http://web.sfc.keio.ac.jp/~rdv/keio/sfc/teaching/architecture/computer-architecture-2013/lec09-smp.html>)

PROGRAMMING FOR DISTRIBUTED SYSTEMS

- ▷ Low-level programming:
 - Socket programming
- ▷ Need to write code to send/receive sockets (messages) through network
 - ▷ Initialize socket in each processor
 - ▷ Sender send message “sendto()”
 - ▷ Receiver get message “recvfrom()”
- ▷ Use this only when you need very low level control



(Figure from [https://people.eecs²³.berkeley.edu/~culler/WEI/lab5/lab5_sockets/sockets_intro.html](https://people.eecs.berkeley.edu/~culler/WEI/lab5/lab5_sockets/sockets_intro.html))

MESSAGE PASSING INTERFACE (MPI)


- ▷ An easier (and standard) way to pass messages in distributed systems
- ▷ C, python, ...
- ▷ Several types of “Collective Communication Routines”
 - ▷ Broadcast
 - ▷ Reduce
 - ▷ Gather, Allgather
 - ▷ ...
- ▷ Check <http://materials.jeremybejarano.com/MPIwithPython/>

```
import numpy
from mpi4py import MPI
comm = MPI.COMM_WORLD

rank = comm.Get_rank()
rankF = numpy.array(float(rank))
total = numpy.zeros(1)
comm.Reduce(rankF, total, op=MPI.MAX)
```



MESSAGE PASSING INTERFACE (MPI)

- ▷ MPI_Broadcast: Broadcasts a message to all other processes of that group



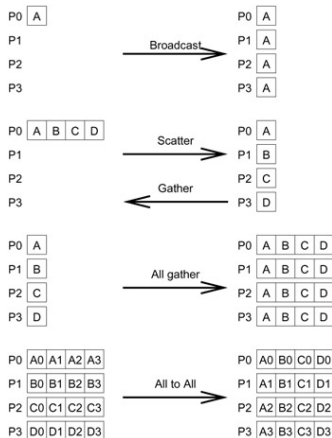
MESSAGE PASSING INTERFACE (MPI)

- ▷ MPI_Reduce: Reduces values on all processes to a single value
(Can specify an operator, e.g., $+$, $-$, \times , $/$)



MESSAGE PASSING INTERFACE (MPI)

▷ Many other operations.



MAPREDUCE PARADIGM

- ▷ Framework for parallel computing
- ▷ Handles parallelization, data distribution, load balancing & fault tolerance
- ▷ Allows once to process huge amounts of data (terabytes & petabytes) on thousands of processors
- ▷ Google
 - Original implementation
- ▷ Apache Hadoop MapReduce
 - Most common (open-source) implementation
 - Built to specs defined by Google
- ▷ Amazon MapReduce
 - On Amazon EC2

MAPREDUCE PARADIGM

- ▷ Framework for parallel computing
- ▷ Handles parallelization, data distribution, load balancing & fault tolerance
- ▷ Allows once to process huge amounts of data (terabytes & petabytes) on thousands of processors
- ▷ Google
 - Original implementation
- ▷ Apache Hadoop MapReduce
 - Most common (open-source) implementation
 - Built to specs defined by Google
- ▷ Amazon MapReduce
 - On Amazon EC2

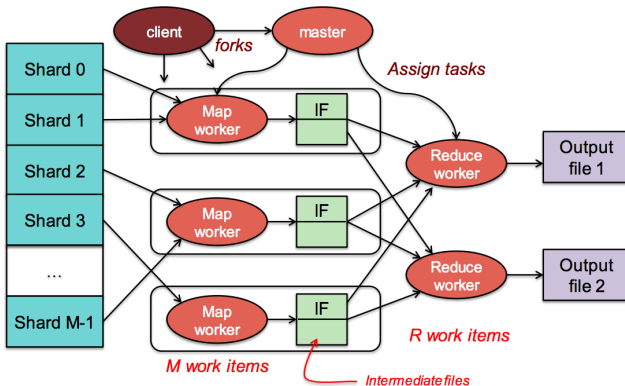
MAPREDUCE CONCEPT

- ▷ Map
 - ▷ Grab the relevant data from the source
 - ▷ User function gets called for each chunk of input (key, value) pairs
- ▷ Reduce
 - ▷ Aggregate the results
 - ▷ Gets called for each unique key

MAPREDUCE CONCEPT

- ▷ Map: (input shard) \rightarrow (intermediate key, intermediate value)
 - ▷ Automatically partition input data to each computer
 - ▷ Group together all intermediate values with the same intermediate key
 - ▷ Pass to the reduce function
- ▷ Reduce: (intermediate key, intermediate value) \rightarrow result files
 - ▷ Input: key, and a list of values
 - ▷ Merge these values together to form a smaller set of values
 - ▷ Automatically partition the reducers distributedly

MAPREDUCE: THE COMPLETE PICTURE



(Figure from <https://www.cs.rutgers.edu/~pxk/417/notes/content/16-mapreduce-slides.pdf>)

EXAMPLE

- ▷ Count number of each word in a collection of documents
- ▷ Map: parse data, output each word with a count (1)
- ▷ Reduce: sum together counts for each key (word)
- ▷ Mapper:

```
map(key, value):  
  // key: document name, value: document contents  
  for each w in value:  
    output (w, '1')
```

- ▷ Reducer:

```
reduce(key, values):  
  // key: a word; values: a list of counts  
  for each v in values:  
    result += Int(v)  
  output(result)
```

EXAMPLE

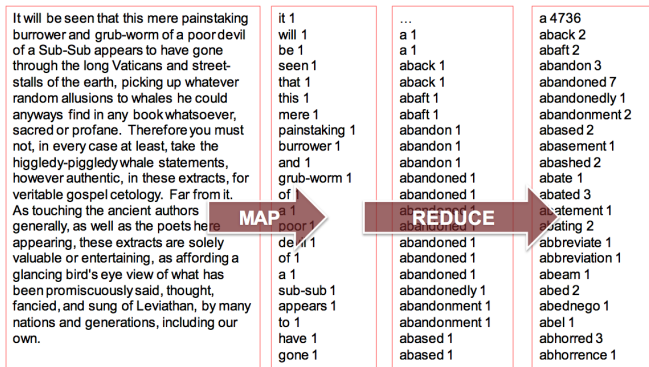
▷ Mapper:

```
for line in sys.stdin:
    line = line.strip().split()
    for word in words:
        print '%s\t%s'%(word,'1')
```

▷ Reducer:

```
word2count = {}
for line in sys.stdin:
    line = line.strip()
    word, count = line.split('\t', 1)
    word2count[word] = word2count[word]+count
for word in word2count.keys():
    print '%s\t%s'%(word,word2count[word])
```

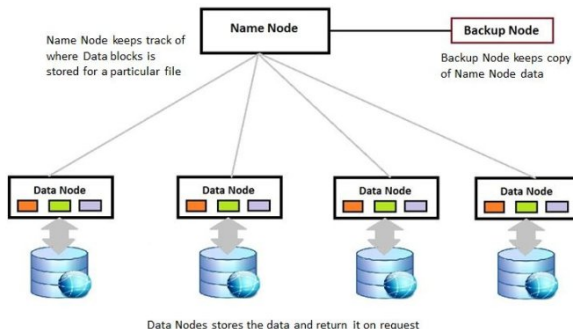
EXAMPLE



(Figure from <https://www.cs.rutgers.edu/~pxk/417/notes/content/16-mapreduce-slides.pdf>)

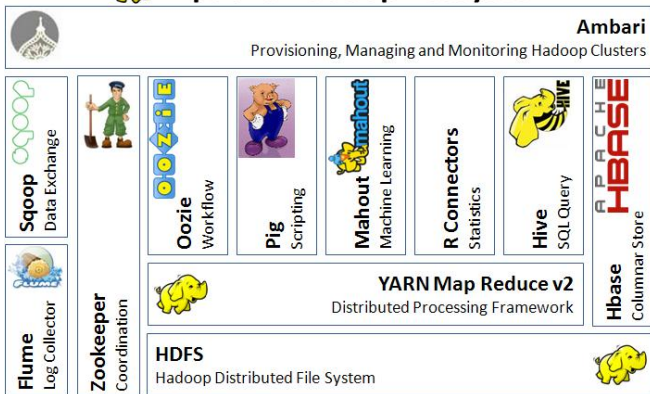
HADOOP DISTRIBUTED FILE SYSTEM (HDFS)

- ▷ The Hadoop Distributed File System (HDFS) is designed to store very large data sets on multiple servers.
- ▷ To use hadoop MapReduce, input/output files are in HDFS







Apache Hadoop Ecosystem



- ▷ Hadoop: Need to read/write HDFS for all the mapper/reducer
Main bottleneck is disk reading time
Not suitable for machine learning (iterative computing)
- ▷ Spark: Also a MapReduce framework with
In memory data flow
Optimize for multi-stage jobs



- ▷ Machine Learning using Spark: [MLlib](#)



PARAMETER SERVER

- ▷ A concept mainly for parallelizing machine learning algorithms (deep learning)
- ▷ Maintain a set of “shared parameters”
- ▷ Local machine communicate with **parameter server** to get the latest parameters

